



# Optimizing Performance Under Power Constraints

Ahmet Inci | NeurIPS | December 2<sup>nd</sup> 2025



NEURAL INFORMATION  
PROCESSING SYSTEMS

# Agenda

- **Rethinking fundamental design metrics in the age of energy and power limits**
  - Interplay of performance, power, and energy
  - How power limits translate to application performance
- **How workload behavior affects performance under power**
  - Workload behavior and sustained performance
  - Impact of data distribution on performance
  - How short-timescale vs. long-timescale power limits shape performance
- **Optimizing performance under power constraints**
  - Performance (Max) vs. Energy (Sum)
  - SW optimization considerations
    - Case study: Energy-aware kernel selection
    - Case study: Utilization-only optimizations
  - Practical optimization guidelines
  - Measuring LLM Inference Efficiency in TRT-LLM
    - Why efficiency matter in real-deployments
    - GPU energy profiling in TRT-LLM

# System Design: Performance (Ethereal) & Energy (Physical)



## Performance (The Goal)

An abstract, emergent metric of **how well** a system functions. It's the conceptual result.

**Metrics:** Speed (Latency), Throughput



## Energy (The Cost)

The **tangible, physical resource cost** of all operations. It's the measurable input as rate of Energy consumption

**Metrics:** Joules (J), Watts (W)



## Efficiency (The Bridge)



Efficiency

$$= \frac{\text{Performance (Useful Work/sec)}}{\text{Energy/sec Consumed}}$$

Connects the abstract goal to the physical cost. This is the key system **trade-off**.

## Performance and Power are Interdependent


Computing consumes energy; rate of computing (performance) consumes power



\*Clock frequency = How many cycles GPU runs per second

"Constraints" driven by: Power Delivery, Thermal, Reliability, etc.

## Performance, Energy and Power



Consider what happens when we optimize SW to reduce execution time?  
Case A: Reduce total work done i.e. # {ops, bytes} (e.g. Kernel fusion)  
Case B: Do not reduce work (e.g. Overlap compute & communication)

\*Simplified view with fixed frequency, voltage etc.



# **Workload Behavior Under Power Constraints**

# Workload Characteristics and Sustained Performance

Why different workloads land at different operating frequencies?

- **Power Budget**

- **TGP (Total Graphics Power)** = total allowed power for GPU
- Enforced and shared across *all subsystems*

- **Power Behavior**

- Power draw depends on workload and clock frequency
- For long-running workloads, average power must stay within TGP
- GPU may exceed TGP briefly (e.g. up to 1820 W on GB300)


- **Workload Performance**

- Sustained frequency differs across workloads
- Dynamic power  $P \approx \text{utilization} \times \text{frequency}$ 
  - Utilization differs across workloads:
    - MatMul - high tensor core util.  $\rightarrow$  high power
    - MemRead – low util.  $\rightarrow$  low power

- **Key insights**

- Compute-bound workloads throttle earlier to stay within power budget
- Memory-bound workloads run at higher clocks because their low compute utilization can't saturate TGP
- Utilization drives power, and power dictates sustained frequency


Different workloads land at different operating points



# Data Distribution Affects Performance

Input sparsity changes switching activity and sustained performance


- Even with the same MatMul kernel, data distribution matters
  - Zeros cause fewer bit flips, so dynamic power is lower
  - With random data, switching activity increases, dynamic power rises and GPU throttles frequency to stay within TGP
- Key insights**
  - Always profile with representative data distributions
    - Zero data underestimates power and overestimates sustained clocks and throughput



# Power Constraints Across Timescales

Short vs. long timescale power limits and how they shape workload performance?

- Power constraints aren't a single number (e.g. GB300 TGP = 1400 W)
  - Power = rate of energy consumption (Joules / sec)
  - Limits are enforced as moving averages, not instantaneous caps
    - Constraints define how fast we're allowed to spend that energy
- Short timescale (Electrical Design Point: 100s of us to 10s of ms)
  - Peak transient power over a short moving average window
- **Key insights**
  - Short timescale limits shape kernel-level performance
  - Compute intensive kernels may momentarily run at higher frequency
    - Brief boost allowed, but frequency is clamped if short-timescale energy budget is exceeded





Short timescale limits shape instantaneous kernel behavior;  
Long timescale limits determine where workloads settle over time

# Workload Transitions and Perf under Power

How GPUs adapt to workload compute-intensity

- GPU power controller operates over a moving average power window
  - GPU adjusts frequency to keep average power  $\leq$  TGP, not instantaneous power
- Short MatMul can draw high power as long as the moving average power  $\leq$  TGP;
  - Power controller doesn't react – GPU sustains its frequency
- Each transformer block cycles through these high-power (MatMul) and low-power (non-MatMul) phases
  - End-to-end performance is shaped by workload composition, not individual kernel clocks




High MatMul activity pushes average power to TGP  $\rightarrow$  GPU lowers frequency to stay within TGP

Non-MatMul phases reduce average power  $\rightarrow$  GPU allows higher frequency

Low overall activity  $\rightarrow$  average power stays well below TGP  $\rightarrow$  GPU sustains near  $\sim F_{max}$

## How can we study this?

Interleaved MatMul and MemRead micro-benchmarking to mimic end-to-end application behavior



\*Speedups vary based on TGP/EDPp ratios, based on GPU

### Micro-benchmark details

- **MatMul** = highly optimized MatMul kernel (>90% util, random data)
- **MemRead** = kernel reading from memory (~90% util, random data)
- Kernels are interleaved
- MatMul duty cycle = % of time in MatMul



### Observations

- 100% MatMul is the most power-limited case
  - Worst-case delivered FLOPS at TGP
  - Pessimistic since most workloads have non-MatMul phases
- Interleaving MemRead increases MatMul performance
  - Low-power phases (MemRead) let MatMul run faster in short bursts
    - During MatMul, power exceeds TGP but gets averaged out to TGP eventually
  - Diminishing returns as duty cycle decreases
    - Average power must satisfy short-timescale limits

Real workloads outperform pure MatMul microbenchmarks because low-power phases reduce average power and allow higher performance for MatMuls

# Transformer Block as an Example

Power varies across kernels, but GPU power controller sees only the averaged behavior



DeepSeek-v3 Transformer Block

Fluctuations at small timescales get averaged out over a transformer block




# **Optimizing Performance Under Power Constraints**

# Performance (Max) vs. Energy (Sum)


Performance depends on the slowest component; Energy depends on all components

- Performance =  $\max f(x)$ 
  - End-to-end performance is determined by the critical path
  - Speeding up non-critical work does not improve performance
  - Slowest component at any level {warp -> block -> kernel -> layer} determine
    - $T_{\text{total}} = \max (T_1, T_2, \dots, T_N)$



- Energy =  $\sum f(x)$ 
  - Energy is accumulated over all kernels
  - Reducing energy in any kernel improves overall efficiency
    - $E_{\text{total}} = \sum_i P_i \times t_i$

- Max vs. Sum applies at every level
  - Within kernels -> across kernels -> across layers -> across the model



Performance is gated by the max; energy is driven by the sum  
This asymmetry is why energy optimization is always useful, but performance optimization is selective

## SW Optimization Considerations

When SW reduces actual work, the GPU runs faster and more efficiently

- Case Study: MatMul Tile Shapes
  - Default kernel (K1): **128x128** tile
  - Alternate kernel (K2): **128x256** tile

MNK:1024x128000x1024 measured on H100-SXM	K1	K2	Delta
Runtime @ 1400 MHz (us)	<b>277.9</b>	291	-4.5%
Power @ 1400 MHz (W)	598	510	17.3%
Runtime @ 700 W (us)	<b>267.2</b>	<b>267.4</b>	-0.07%
Frequency @ 700 W (MHz)	1462	<b>1530</b>	-4.5%
Runtime @ 400 W (us)	370.1	<b>335.8</b>	10.2%
Frequency @ 400 W (MHz)	1045	<b>1215</b>	-14%

At same frequency, K1 consumes more power and is faster

At fixed power of 700 W, runtime difference is negligible


At a lower power limit, K2 is ~10% faster

- Why does K2 behave differently?
  - Larger tiles improve data reuse
  - Fewer L2 bytes → lower energy → faster at lower power (perf/W ↑)

## What happens when you optimize a piece of SW?

Increasing utilization without reducing work has limited benefit

- Baseline
  - Compute and communication run sequentially
  - Power swings between high (MatMul) and low (All-Reduce)
  - ~60-70% Math utilization, runs at frequency =  $F_1$
- After SW optimization (Compute-Communication Overlap)
  - Overlapping compute-communication increases utilization
  - Combined activity increases average power
  - What happens to overall performance?
    - Performance improves, but less than expected utilization gain
    - Higher average power -> lower sustained frequency  $F_2 < F_1$
  - Did the optimization reduce total work?



True performance gains come from reducing work, not just keeping the GPU busy


# Optimization Guidelines

- Simple optimization guidelines:
  - Objective 1: Reduce cycles (increase utilization)
  - Objective 2: Reduce work (fewer ops, fewer bytes -> lower power)
- This is business as usual for GPU developers:
  - Understand and overcome throughput bottlenecks (blocking, data reuse, etc.)
  - Hide latencies (software pipelining, etc.)
  - Overlap compute and non-compute work
  - Improve algorithms to minimize work
- Optimizations that reduce cycles usually also reduce energy
  - Things that hurt performance also hurt energy – wasted ops, extra memory traffic, etc.
    - Example: tuning MatMul traversal order to increase L2 hit rates, reduces DRAM bytes fetched, saves cycles and energy
    - Counter-examples: speculation (with high misprediction rates); operations that are off the critical path

# You Can't Optimize What You Can't Measure: Energy Monitoring in TRT-LLM

## Why tokens/sec/Watt matter for real-world deployment?

- First step to optimize: Measure the cost of inference
  - Efficiency = Tokens per Joule = Tokens/sec/Watt
- Understand your use case:
  - What is the cost of interactivity?
    - Low-latency scenarios sacrifice efficiency
    - High-throughput scenarios provide best Perf/W
  - How much efficiency are you willing to trade for responsiveness?
- Datacenters operate under fixed power budgets
  - Perf/W determines the "factory" throughput in DC power budget
    - Higher Perf/W -> more GPUs -> higher overall TPS



Running GPUs at their best Perf/W point enables more GPUs within the same DC power budget and higher overall throughput

# GPU Energy Profiling in TRT-LLM

How to measure GPU energy and identify efficiency improvements?

## Goals

- Enable on-the-fly power and energy tracking in TRT-LLM
- Identify Pareto-optimal configs for performance and energy

## How to Use

- Collect energy data using public NVML APIs:
  - **GetTotalEnergyConsumption**
    - reports energy consumed since last driver call
  - Automatically logs energy if **pynvml** is available
    - default in most NGC containers
  - Logged with other performance metrics:

```
-- Energy Metrics --  
Total Energy (J): 160237.5460  
Output Tokens per Joule (tokens/J): 13.0878  
Average GPU Power (W): 991.4987
```

Shout out to ML.Energy team for their related work in energy profiling!





**Thank you!**

**We're Hiring!**  
Help Us Shape the Future





# **Backup**

# Sample Commands

How to use *nvidia-smi*

- **Measure**

- Sample nvidia-smi command:
  - `nvidia-smi --query-gpu=timestamp, clocks.sm, power.draw, power.draw.instant, temperature.gpu --format=csv -lms 100`

- **Control**



- Change power limit
  - `sudo nvidia-smi -pl 400`
- Adjust on-chip clock ratios
  - `sudo nvidia-smi boost-slider --vboost {0, 1, 2, ..}`
- Lock clocks
  - `sudo nvidia-smi -lgc 1500`

- To learn more <https://docs.nvidia.com/deploy/nvidia-smi>

# Power Stabilization for AI Datacenters

## In the Age of Gigawatt Datacenters

- AI workloads cause large power fluctuations in datacenters
  - ~1000s of GPUs consume power and idle simultaneously
  - Harder for datacenters to provision power
  - Leads to throttling and lower total throughput
- Synchronous nature of training workloads
  - High power / compute-heavy phase
    - GPUs work on local data (e.g. MatMul)
      - Tensor cores operate at high utilization, drawing power close to TGP
  - Low power / communication phase
    - GPUs sync on the data (e.g. All-Reduce, checkpointing)
      - Utilization drops sharply, power collapses




Unlike traditional DC jobs, in AI workloads GPUs operate in lockstep,  
~1000s of GPUs breathing power in and out

# Steady Power Vision

Stable rack-level power improves sustained ML performance

- Energy buffering smooths bursty ML workload power
  - Charge during low-utilization phases (communication, gradient sync)
  - Discharge during high power bursts (MatMul-heavy compute)
  - Keeps rack-level power stable even when GPU activity is spiky
- **Why should ML practitioners care?**
  - Enables provisioning for average, not peak power
    - More GPUs can fit within the same datacenter power budget
    - Reduces throttling and improves overall TPS/Watt
- **How can SW support steady power?**
  - Design models with balanced compute and communication
    - Avoid big power swings from back-to-back MatMul or All-Reduce
    - Overlap compute and communication
  - Profile per layer performance and power
    - Identify kernels that cause transient peaks
  - Favor scheduling strategies that flatten utilization across GPUs
    - Interleaved pipeline scheduling (1F1B, ZeroBubble, DualPipe)



<https://developer.nvidia.com/blog/how-new-qb300-nvl72-features-provide-steady-power-for-ai/>

Stabilizing power at the DC is not just an electrical problem, it is a system-level enabler for scaling ML performance

## GB300 PSU Measured Benefits

Energy buffering reduces peak power, allowing racks to stay within provisioned limits

